Note: *I recommend opening this side by side with the original articles, which are linked below. Then use this document to indicate which parts to read. You may also just read this document if that works better for your reading comprehension.*
The original pages have better formatting, which also makes the footnotes/citations easy to mouse-over and look at (instead of being formatted as numbers like this[1]).

From [Executive Summary of "Artificial Power: 2025 Landscape Report"](#) AI Now Institute (April 2025):

Executive Summary

… [Dr. Greg: omitted for overlap with the AI Now 2023 report]

The good news is that the road offered by the tech industry is not the only one available to us. This report explains why the fight against the industry's vision for AI is a fight worth having, even as we turn ourselves tirelessly toward the task of building out the shared project of a just, equitable, sustainable, and democratic society.

Over the past decade, taming the power of big technology platforms like Microsoft, Amazon, Google, and Meta has increasingly become a central question in American political and public discourse. Unless we contend with the power vested in these firms, we won't meaningfully be able to hold the industry accountable to the interests of the broader public, even as these companies reshape markets, institutions, and infrastructures core to public life.

What does AI have to do with any of this? As we argued in our 2023 report,[1] AI is fundamentally about concentration of power in the hands of Big Tech. At the start of the year, it seemed like the market was poised for disruption, with a new crop of Silicon Valley challengers gaining prominence, like OpenAI, Anthropic, StabilityAI, and Inflection AI. But now, just two years later, it is clear that the bench of key players in this market hasn't changed much: Microsoft, Google, Meta, Musk's xAI, OpenAI (backed by Microsoft), and Anthropic (backed by Amazon and Google).2 The new suite of LLM-powered AI products has pushed these firms into the spotlight, dominating headlines and, increasingly, becoming the subject of dinner-table conversation.

Amid the frenzy, there's been a misplaced focus on blinkered questions of whether one AI system or application is good or bad, or evaluating the moral quandaries of hypothetical worlds. Instead we need to redirect attention to the AI ecosystem, and its dependencies and risks, as a whole. **The question we should be asking is not if ChatGPT is useful or not, but if OpenAI's unaccountable power, linked to Microsoft's monopoly and the business model of the tech economy, is good for society.**

… [Dr. Greg: omitted for length]

But it's not just market power we need to be concerned with: These tech oligarchs are counting on a wholesale rewriting of our social and economic foundations, using AI as the justification. From breaking apart the US federal government and raiding citizen data under the guise of efficiency, to redesigning workflows to devalue human labor and creativity so they are AI-ready, to redirecting our entire energy infrastructure to prioritize their technology over people's basic needs, the vision promulgated by tech oligarchs requires, as a foundation, the unraveling of core social, political, and economic fabrics.

… [

Dr. Greg: omitted for length. Argues existing deployed AI systemsconcentrate power and reduce transparency, even before the LLM / ChatGPT era.

For example, automated hiring systems screen and reject applicants, and applicants have no idea when or why they are rejected by AI in the early stages. The lack of such features is not accidental, it is a design decision made by companies that build those systems, and laws/regulations could be passed to mandate such transparency features, or even mandate evaluating bias/fairness in such systems.

Laws/regulations could require demonstrating these AI systems

perform better than previous manual screening before allowing the AI system to be deployed (i.e. pre-market approval).
]

Why society would ever accept this bargain is the critical question at hand. Amid the excitement over AI's (speculative) potential, the sobering reality of its present and recent past is obscured. When we consult the record on how AI is already intermediating critical social infrastructures, we see that it is materially reshaping our institutions in ways that ratchet up inequality, render institutions opaque to those they are meant to serve, and concentrate power in the hands of the already powerful. (See **Chapter 3: Consulting the Record**.) It makes clear that for all the whiz-bang demos and bold Davos proclamations, on the ground AI is consistently deployed in ways that make everyday people's lives, material conditions, and access to opportunities worse and the systems that incorporate them stronger.

This report's title, *Artificial Power,* captures the critical, and at times contradictory, moment we find ourselves in. On one hand, the tech oligarchy has successfully deployed "AI"—as a strategic marketing term and as a set of automation technologies—to cement and grow its power. At the same time, this power is vastly inflated, contingent, and poised for disruption. Contending with the dual reality of how those with power have deployed AI systems to enact significant harm while exposing the ways this power can and must be disrupted is the central work of this moment.

The Elephants in the AI Room:

Of Business Model(s) and Fatal Technical Flaws

The AI industry is on shakier ground than it may seem. Valuations are sky high, while the business model hinges on an intensely expensive technology that lacks a consistent revenue stream. AI companies bleed money for every user they gain: Anthropic burned through $5.6 billion[6] this year but was valued at $61.5 billion.[7] OpenAI lost $5 billion[8] but is valued at $300 billion[9]. No profit-making use cases exist yet, or are even on the horizon. This may seem like business as usual for the move-fast-and-break-things ethos of Big Tech, but we are in a profoundly different monetary environment now—it's no longer the 2000s or 2010s.

Markets are saturated, market dominance has been established among the platform and infrastructure winners that emerged from those decades, and, put simply, the cost of large-scale AI is eye-watering at a level not seen before in tech.

The question now circulating more, and more openly, is this: When will the AI bubble burst and who will be impacted by it? Because this *is* a bubble. For their part, companies are pushing out shiny objects to detract from the business reality while they desperately try to derisk their portfolios through government subsidies and steady public-sector (often carceral or military) contracts. While it's very clear how tech companies benefit from claims to "public-interest AI" used to justify the pouring of taxpayer dollars into this industry, it is not at all clear how this benefits the rest of society (See **Chapter 2: Heads I Win, Tails You Lose**).

Chapter Snapshot

Heads I Win, Tails You Lose

This section maps the drivers that are securing Big Tech firms' advantage in the AI market, before turning to the question of who loses in the end:

- Cloud infrastructure providers benefit from cycles of AI dependence
- Big Tech firms benefit from leveraging control over the tech ecosystem
- Big Tech benefits from the data center boom … even if the AI boom doesn't pan out

… [Dr. Greg: omitted for length, important points are made in longer paragraphs throughout but are omitted for a first read to focus on the overall summary / snapshots of each chapter of the report from this point on]

Chapter Snapshot

Consulting the Record

This section compiles over a decade of evidence showing how the tech industry has sought to reshape society to enable more widespread deployment of the technologies it builds and profits from, often contributing to the degradation of our social, political, and economic lives. The section aligns on five key takeaways:

- **AI's benefits are overstated and underproven**, from cancer cures to hypothetical economic growth – while some of its flaws are real, immediate, and growing.
- **AI-sized solutions to entrenched social problems displace grounded expertise**, in disparate domains like higher education, healthcare, and agriculture.
- **AI solutionism obscures systemic issues facing our economy – obscuring economic concentration and acting as a conduit for deploying austerity mandates by another name**. The DOGE power grab is instructive, though New York's MyCity offers another example where millions of taxpayer dollars were invested into flawed AI solutions that failed to deliver tangible benefits to the public.
- **The productivity myth obscures a foundational truth – the benefits of AI accrue to companies, not to workers or the public at large**, even as algorithmic management tools make work unstable and unsafe. 'Agentic AI' will make workplaces even more bureaucratic and surveillant, reducing not increasing autonomy.
- **AI use is frequently coercive, violating rights and undermining due process.** This is nowhere more clear than the rise of AI usage in immigration enforcement, where human rights abuses are common and legal norms are routinely violated – even before AI is in the mix.

Chapter Snapshot

AI's False Gods

This section interrogates narratives that advance AI industry dominance and make the current trajectory of AI seem inevitable:

- **The AGI Mythology: The Argument to End All Arguments** unpacks the nebulous claims surrounding "artificial general intelligence," arguing that the term collapses complex technical realities into a singular, imminent, and inevitable future that conveniently advances the interests of the companies claiming to build it.

- **'Too Big to Fail:' Infrastructure and Capital Push** explores how tech firms are deploying unprecedented amounts of capital to perpetuate a "bigger-is-better" AI paradigm, shoring up their continued market dominance through government and taxpayer support.
- **AI Arms Race 2.0: From Deregulation to Industrial Policy** details how the US-China AI arms race has heightened, and is now used to brand a slate of industrial policy initiatives designed to boost the tech industry and avert regulatory scrutiny.
- **Recasting Regulation as a Barrier to Innovation** shows how the AI industry has strategically pitted regulation against innovation, leading to a global deregulatory posture that ignores the role regulation plays in enhancing innovation and competition.

Chapter Snapshot

A Roadmap for Action: Make AI a Fight About Power, Not Progress.

1. **Target how the AI industry works against the interests of everyday people.** In the aftermath of the 2024 election, there's growing consensus across the ideological spectrum that focusing on the material conditions and economic interests of working people is key to building political power. We need to not only make AI-related issues more relevant to movements fighting for economic populism and against tech oligarchy; we also need to better target the AI industry as a key actor working *against* the interests of the working public. The pushback against the Department of Government Efficiency (DOGE), the buildout of AI data centers, and algorithmic prices and wages constitutes fertile ground for building a broader movement unified in its focus on rejecting AI's unaccountable tech-enabled social and political control.
2. **Advancing worker organizing is the clearest path to protecting us and our institutions from AI-enabled capture.** Labor campaigns have demonstrated that working people have a particular form of power to wield, power that can determine how their employers deploy AI and digital systems. The deeper opportunity for labor, and a more transformative ambition,

however, would be to direct labor's power not just toward whether and how AI is used in the workplace, but also toward recalibrating the technology sector's power overall and shaping the trajectory of AI in the public interest and common good.

3. **Enact a 'zero trust' policy agenda for AI.** Trust in AI firms' benevolence is not a smart, informed, or credible option—not if we're going to proceed with serious work. Enacting a policy agenda built on bright-line rules that restrict the most harmful uses of AI, regulate the AI lifecycle from nose to tail, and ensure that the industry that currently creates and profits from AI isn't left to regulate and evaluate itself—essentially grading its own homework—must be a priority at the state and federal levels in the US, and internationally.
4. **Bridge networks of expertise, policy, and narrative to strengthen AI advocacy.** AI advocacy and policy has often been undermined too often by blinkered views that fail to see the different components of the AI supply chain materially, are often single issue-focused, and it is easy to miss the ways in which big picture narratives manifest to limit possibilities in policy fights. From national security logics that can be a vector both for, and cutting against, moves towards industry accountability; to reframing traditional data privacy levers as key tools in the fight against automation and addressing market power.
5. **Reclaim a positive agenda for public-centered innovation without AI at the center.** The current trajectory of AI puts the public under the heel of unaccountable tech oligarchs. But their success is not inevitable. By moving out from under the shadow of the idea that large-scale AI is inevitable, we can reclaim the space to conduct real innovation and to push for exciting and novel alternative pathways that leverage tech to shape a world that serves the public and is governed by our collective will.